

RETRIEVAL WITHOUT GENERATING AN ANSWER

Can a decision model search a whole wiki?

We gave Jev a real 216-page company wiki and a simple job: explore the tree, follow page links, and return up to ten pages worth reading.

2.00s

median query latency

\$0.00134

mean gateway cost /
query

10 / 11


reference pages found

36

runs · 12 questions × 3

Promising for natural-language retrieval. Jev matched the reference-page recall of hand-chosen keywords while accepting a detailed question directly. This small experiment does not establish better recall than keyword search, exhaustive coverage, or an LLM speed advantage.


Better questions in. Comparable recall out.



11 preselected reference pages across 10 answerable questions. Jev found the same 10 references in all three repetitions. Search baselines ran once each against the same live wiki; keywords were manually chosen. This is reference-page recall, not full-corpus recall or precision.

One multi-part question consistently missed a secondary reference. Two deliberately implausible questions returned no matches in all six runs. Those are easy negative cases; realistic near-misses remain untested.

A two-second median. A longer tail.



Median **2.00s** · p95 **5.21s**.

Timing includes sequential gateway decisions and traversal over a local snapshot. It excludes production database, parent-agent, and connector latency.

21 of 36 runs hit a coverage limit.

Limits include branch pruning, page excerpts, page reads, and link rounds. Bounded retrieval makes latency manageable; it also prevents a claim of exhaustive search.

The model decides. The harness does the work.

[Detailed question](#)[Score tree entries](#)[Read candidates](#)[Follow wiki links](#)[Return 0-10 matches](#)

Jev evaluates typed Boolean questions in parallel within each request. The harness turns those scores into folder traversal and page reads. Candidate titles select where to look; page content determines which results survive. The parent session retrieves the matches to construct its answer.

A small, bounded loop

20 seconds · 20 model calls · 48 pages ·
12,000 characters per page · at most eight
folders expanded per level. Navigation
threshold 0.35; final-match threshold 0.70.
Scores are uncalibrated for this task.

Source retrieval is still a separate step

Internal wiki links are followed. External source pointers are returned to the main session for authorized retrieval. This experiment does not measure fetching Slack, email, issue, or web source bodies.

Candidate identity changed the result.

The first 12-query pilot found **zero of the 11 reference pages**. Evaluation questions pointed to candidates only by array index. In a two-page diagnostic, both pages received 0.89 relevance even though only one matched the question.

Adding each candidate's explicit path and title to its evaluation question changed that diagnostic to **0.89 versus 0.03**. The revised harness then found 10/11 references consistently. This suggests candidate binding matters, but one prompt change and one diagnostic do not establish a universal Jev limitation.

Inspect a measured result.

Replay of recorded results, not a live model call. Private page titles, bodies, source references, and full questions are omitted. Page IDs are stable only within this report.

Question category	References found	Median time	Mean cost
Q01 · CI bottleneck	1 / 1	1.13s	\$0.00075
Q02 · Durable reports	1 / 2	1.74s	\$0.00105
Q03 · User input interruption	1 / 1	3.04s	\$0.00171
Q04 · Sync migration	1 / 1	1.24s	\$0.00062
Q05 · Pricing comparison	1 / 1	4.73s	\$0.00307
Q06 · Video retention	1 / 1	2.75s	\$0.00121
Q07 · Back navigation	1 / 1	2.50s	\$0.00138
Q08 · Feature retirement	1 / 1	1.80s	\$0.00096
Q09 · Local inference	1 / 1	1.92s	\$0.00103
Q10 · Agent evaluations	1 / 1	4.50s	\$0.00326
Q11 · No answer A	No answer	1.62s	\$0.00095
Q12 · No answer B	No answer	0.53s	\$0.00014

About five cents to learn something useful.

The bill

The 36 benchmark runs cost **\$0.04836**. The pilot, three diagnostic calls, and two API integration checks bring measured gateway spend to **\$0.05275** of the authorized \$5 budget. Costs come from gateway response metadata and include its reported surcharge; this is not a reconciled invoice.

The dataset

216 pages, 22 folders, 793,066 body characters. One workspace's default wiki, captured on September 17, 2026. All pages were available to the harness; each query read only its selected subset. No private corpus content is included in this report.

Method: Ten task questions plus two easy unanswerable questions; 11 reference pages selected from inspected content before the scored runs. One pilot, one prompt correction, then three sequential repetitions with the same settings. p95 uses nearest rank. No embedding index, answer generation, LLM baseline, blind human labeling, full relevance judgments, or held-out corpus. Repeated queries may benefit from provider caching. Latency should not be compared directly with the live search baseline's transport time.

Next test: A larger, independently labeled set with realistic near-misses, branch ambiguity, stale facts, and evidence deep inside long pages. Measure end-to-end production latency and precision before claiming search quality or speed superiority.

References: [Vercel model and pricing page](#) · [Vercel Jev announcement](#). Model: typesafe-ai/jev ; AI SDK 7.0.105; zero data retention requested; retries disabled.